

# Data Analyst Interview Questions (0-3 Years) 17-19 Ipa

## SQL Questions

### 1. Write a query to find duplicate rows in a table.

To detect duplicates, identify columns that should be unique and group by them.

#### Example:

```
SELECT column1, column2, COUNT(*) AS count
FROM your_table
GROUP BY column1, column2
HAVING COUNT(*) > 1;
```

#### Explanation:

- GROUP BY combines rows with the same values in the specified columns.
- HAVING COUNT(\*) > 1 filters those combinations that occur more than once, indicating duplicates.

\* Tip: Add ROW\_NUMBER() or RANK() with CTE to highlight or delete duplicates if needed.

### 2. Explain the difference between INNER JOIN and OUTER JOIN with examples.

#### ◆ INNER JOIN:

Returns only **matching** records from both tables.

```
SELECT e.name, d.department_name
FROM employees e
INNER JOIN departments d ON e.department_id = d.department_id;
```

- Output: Only employees who belong to a department.

#### ◆ LEFT OUTER JOIN:

Returns **all records from the left** table, and matching records from the right table. If no match, NULL is returned.

```
SELECT e.name, d.department_name
FROM employees e
LEFT JOIN departments d ON e.department_id = d.department_id;
```

- Output: All employees, with department info where available.

#### ◆ RIGHT OUTER JOIN:

Returns **all records from the right** table, and matching records from the left.

#### ◆ FULL OUTER JOIN:

Returns **all records from both tables**, matching where possible.

**Key Difference:**

- INNER JOIN = intersection (matched data only)
- OUTER JOIN = union + NULLs (matched + unmatched data)

### 3. Write a query to fetch the second-highest salary from an employee table.

**Option 1: Using DISTINCT, ORDER BY, and LIMIT (MySQL/PostgreSQL)**

```
SELECT DISTINCT salary  
FROM employees  
ORDER BY salary DESC  
LIMIT 1 OFFSET 1;
```

**Option 2: Using subquery (Generic SQL)**

```
SELECT MAX(salary)  
FROM employees  
WHERE salary < (SELECT MAX(salary) FROM employees);
```

**Explanation:**

- The subquery fetches the highest salary.
- The outer query finds the maximum salary **less than** the highest — giving the second-highest.

### 4. How do you use GROUP BY and HAVING together? Provide an example.

Use GROUP BY to group data and HAVING to filter **aggregated results** (unlike WHERE, which filters raw rows).

```
SELECT department_id, COUNT(*) AS emp_count  
FROM employees  
GROUP BY department_id  
HAVING COUNT(*) > 5;
```

**Explanation:**

- Groups employees by department.
- Filters groups where the count of employees is **more than 5**.

### 5. Write a query to find employees earning more than their managers.

Assume the table employees has:

emp\_id, name, salary, manager\_id

```
SELECT e.name AS employee_name, e.salary, m.name AS manager_name, m.salary AS  
manager_salary  
FROM employees e  
JOIN employees m ON e.manager_id = m.emp_id  
WHERE e.salary > m.salary;
```

**Explanation:**

- Self-join: matches employees (e) with their managers (m).
  - Filters those where employee's salary > manager's salary.
- 

## 6. What is a window function in SQL? Provide examples of ROW\_NUMBER and RANK.

### ◆ Definition:

A **window function** performs calculations **across a set of table rows** related to the current row — without collapsing rows like GROUP BY.

### Syntax:

```
FUNCTION_NAME() OVER (PARTITION BY column ORDER BY column)
```

---

### ◆ Example: ROW\_NUMBER()

Assigns a unique sequential number to each row **within a partition**.

```
SELECT name, department, salary,
```

```
    ROW_NUMBER() OVER (PARTITION BY department ORDER BY salary DESC) AS row_num  
FROM employees;
```

- Each employee within the same department gets a row number based on salary rank (highest first).
- 

### ◆ Example: RANK()

Assigns **the same rank** to rows with **equal values**, but skips the next rank(s).

```
SELECT name, department, salary,
```

```
    RANK() OVER (PARTITION BY department ORDER BY salary DESC) AS rank_num  
FROM employees;
```

- If 2 employees have the same salary, both get rank 1, and the next gets rank 3.
- 

## 7. Write a query to fetch the top 3 performing products based on sales.

Assume table sales\_data has:

```
product_id, product_name, total_sales
```

```
SELECT product_id, product_name, total_sales  
FROM sales_data  
ORDER BY total_sales DESC  
LIMIT 3;
```

### Alternate using RANK() (if ties matter):

```
SELECT product_id, product_name, total_sales  
FROM (  
    SELECT *, RANK() OVER (ORDER BY total_sales DESC) AS rank_num  
    FROM sales_data  
) ranked_sales  
WHERE rank_num <= 3;
```

## 8. Explain the difference between UNION and UNION ALL.

Feature	UNION	UNION ALL
Duplicates	Removes duplicates	Keeps all rows, including duplicates
Performance	Slower (because of sorting)	Faster (no de-duplication)
Use case	When you want distinct rows	When duplicates are meaningful

### Example:

```
SELECT city FROM customers
UNION
SELECT city FROM vendors;
→ Returns a unique list of cities.
SELECT city FROM customers
UNION ALL
SELECT city FROM vendors;
→ Returns all cities, including duplicates.
```

## 9. How do you use a CASE statement in SQL? Provide an example.

◆ CASE lets you write conditional logic in SQL (similar to IF/ELSE).

```
SELECT name, salary,
CASE
    WHEN salary >= 100000 THEN 'High'
    WHEN salary >= 50000 THEN 'Medium'
    ELSE 'Low'
END AS salary_category
FROM employees;
```

### Explanation:

- Assigns a category based on salary value.
- Works inside SELECT, WHERE, ORDER BY, etc.

## 10. Write a query to calculate the cumulative sum of sales.

Assume table sales has:

order\_date, product\_id, sales\_amount

```
SELECT order_date, product_id, sales_amount,
SUM(sales_amount) OVER (PARTITION BY product_id ORDER BY order_date) AS
cumulative_sales
FROM sales;
```

### Explanation:

- SUM(...) OVER (...) calculates a **running total** per product based on order date.
- PARTITION BY groups by product, and ORDER BY ensures the accumulation follows chronological order.

## 11. What is a CTE (Common Table Expression), and how is it used?

### ◆ Definition:

A **CTE (Common Table Expression)** is a temporary, named result set that you can reference within a SQL query. It improves readability and simplifies complex subqueries or recursive logic.

### ◆ Syntax:

```
WITH cte_name AS (
    SELECT ...
)
SELECT * FROM cte_name;
```

### ■ Example – Filter top-paid employees using CTE:

```
WITH HighEarners AS (
    SELECT emp_id, name, salary
    FROM employees
    WHERE salary > 100000
)
SELECT * FROM HighEarners;
```

### Benefits:

- Reusable and readable
- Allows recursion (e.g., hierarchical data)
- Avoids repeating subqueries

## 12. Write a query to identify customers who have made transactions above \$5,000 multiple times.

Assume transactions table has:

customer\_id, transaction\_amount

```
SELECT customer_id, COUNT(*) AS high_value_txns
FROM transactions
WHERE transaction_amount > 5000
GROUP BY customer_id
HAVING COUNT(*) > 1;
```

### Explanation:

- Filters high-value transactions (> \$5000).
- Groups them by customer.
- Returns customers who've done this **more than once**.

## 13. Explain the difference between DELETE and TRUNCATE commands.

Feature	DELETE	TRUNCATE
Removes rows WHERE supported?	Yes (can use WHERE condition)  ■ Yes	Yes (removes all rows)  + No
Logging	Logs each deleted row (slower)	Minimal logging (faster)
Rollback	■ Can be rolled back (if within transaction)	■ Can be rolled back (in some RDBMS)
Identity reset	+ Retains identity	■ Resets identity (in most DBs)
Use case	Partial deletion or audit trail needed	Full data wipe without audit needed

## 14. How do you optimize SQL queries for better performance?

Here are **key SQL optimization techniques**:

◆ **1. Use SELECT only required columns**

-- Bad

```
SELECT * FROM orders;
```

-- Good

```
SELECT order_id, customer_id FROM orders;
```

◆ **2. Create proper indexes**

- Index frequently used columns in JOIN, WHERE, ORDER BY.

◆ **3. Avoid functions on indexed columns**

-- Slower (cannot use index)

```
WHERE YEAR(order_date) = 2024
```

-- Better

```
WHERE order_date BETWEEN '2024-01-01' AND '2024-12-31'
```

◆ **4. Use EXISTS instead of IN (for subqueries)**

-- Prefer EXISTS (better for large datasets)

```
SELECT name FROM customers c
```

WHERE EXISTS (

```
    SELECT 1 FROM orders o WHERE o.customer_id = c.customer_id
);
```

◆ **5. Avoid unnecessary joins or nested subqueries**

◆ **6. Use appropriate data types and avoid implicit conversions**

◆ **7. Analyze execution plans (EXPLAIN or EXPLAIN ANALYZE)**

## 15. Write a query to find all customers who have not made

## **any purchases in the last 6 months.**

Assume:

- customers(customer\_id, name)
- transactions(customer\_id, transaction\_date)

```
SELECT c.customer_id, c.name
FROM customers c
LEFT JOIN transactions t
ON c.customer_id = t.customer_id
AND t.transaction_date >= CURRENT_DATE - INTERVAL '6 months'
WHERE t.customer_id IS NULL;
```

**Explanation:**

- LEFT JOIN includes all customers.
- WHERE t.customer\_id IS NULL ensures the customer had **no purchase in the last 6 months.**

## **16. How do you handle NULL values in SQL? Provide examples.**

◆ **NULL represents missing or unknown data.**

### **1. Using IS NULL / IS NOT NULL:**

```
SELECT * FROM employees WHERE manager_id IS NULL;
```

### **2. Replace NULL using COALESCE() or IFNULL() (MySQL):**

```
SELECT name, COALESCE(phone_number, 'Not Provided') AS contact
FROM customers;
```

### **3. Handling NULLs in aggregation (e.g., AVG, SUM):**

- These functions **ignore NULLs by default.**

```
SELECT AVG(salary) FROM employees;
```

### **4. Conditional checks:**

```
SELECT name,
CASE
    WHEN salary IS NULL THEN 'Unknown'
    ELSE 'Known'
END AS salary_status
FROM employees;
```

## **17. Write a query to transpose rows into columns.**

Assume a table sales with:  
region, month, sales\_amount

We want to **pivot month values** into columns.

### **Using CASE:**

```

SELECT region,
    SUM(CASE WHEN month = 'Jan' THEN sales_amount ELSE 0 END) AS Jan,
    SUM(CASE WHEN month = 'Feb' THEN sales_amount ELSE 0 END) AS Feb,
    SUM(CASE WHEN month = 'Mar' THEN sales_amount ELSE 0 END) AS Mar
FROM     sales
GROUP BY region;

```

#### ■ Using PIVOT (SQL Server or Oracle syntax):

```

SELECT region, [Jan], [Feb], [Mar]
FROM (
    SELECT region, month, sales_amount
    FROM sales
) AS src
PIVOT (
    SUM(sales_amount)
    FOR month IN ([Jan], [Feb], [Mar])
) AS p;

```

## 18. Explain indexing and how it improves query performance.

#### ◆ What is an index?

An **index** is a data structure that improves the speed of data retrieval operations on a database table at the cost of additional space and write-time performance.

#### ◆ How indexing helps:

Feature	With Index	Without Index
Search performance	Fast (uses binary/tree search)	Slow (scans every row — full scan)
Used in	WHERE, JOIN, ORDER BY, GROUP BY	Inefficient for large datasets
Types	B-tree (default), Bitmap, Hash, etc.	-

#### ■ Example:

-- Creating index  
CREATE INDEX idx\_customer\_id ON transactions(customer\_id);  
• This helps queries like:  
SELECT \* FROM transactions WHERE customer\_id = 101;

#### ! Important notes:

- Too many indexes can slow down INSERT/UPDATE.
- Avoid indexing columns with **low cardinality** (e.g., gender).
- Use **composite indexes** when querying multiple columns together.

## 19. Write a query to fetch the maximum transaction amount for each customer.

Assume a transactions table:

Column	Description
customer_id	ID of the customer
transaction_id	Unique transaction ID
amount	Transaction amount

#### ■ **Query:**

```
SELECT customer_id, MAX(amount) AS max_transaction
FROM transactions
GROUP BY customer_id;
```

#### ❑ **Explanation:**

- GROUP BY groups all transactions by customer.
- MAX(amount) returns the highest transaction for each group (customer).

## 20. What is a self-join, and how is it used?

#### ◆ **Definition:**

A **self-join** is a regular join where a table is joined with itself. It is useful when rows in a table are related to other rows in the same table.

#### ■ **Example Use Case – Employees and Managers:**

Assume:

emp_id	name	manager_id
1	Alice	NULL
2	Bob	1
3	Carol	1
4	David	2

Here, manager\_id refers to emp\_id of another employee.

#### ◆ **Query: Get employee names along with their manager names**

```
SELECT e.name AS employee_name, m.name AS manager_name
FROM employees e
LEFT JOIN employees m
ON e.manager_id = m.emp_id;
```

#### ❑ **Explanation:**

- e is an alias for employees (as employee).
- m is another alias for the same table (as manager).
- The join links an employee to their manager using manager\_id = emp\_id.

## Data Analysis/Scenario-Based Questions

### 21. How would you design a database to store credit card

## transaction data?

To store credit card transaction data, we need to **normalize** the structure while keeping it **scalable, secure, and query-efficient**.

### ■ Suggested Schema Design:

#### 1. Customers Table

customer\_id (PK), name, email, phone, address

#### 2. Cards Table

card\_id (PK), customer\_id (FK), card\_number (masked), card\_type, status, issued\_date

#### 3. Merchants Table

merchant\_id (PK), name, category, location

#### 4. Transactions Table

transaction\_id (PK), card\_id (FK), merchant\_id (FK),  
transaction\_date, amount, currency, status, location

### ☛ Best Practices:

- Mask sensitive fields (like card numbers).
- Store card\_number as encrypted or tokenized.
- Use partitioning on date fields for faster querying.
- Add indexes on card\_id, merchant\_id, transaction\_date.

## 22. Write a query to identify the most profitable regions based on transaction data.

Assume a transactions table:

(transaction\_id, customer\_id, amount, region, transaction\_date)

### ■ Query to find top 3 profitable regions:

```
SELECT region, SUM(amount) AS total_revenue  
FROM transactions  
GROUP BY region  
ORDER BY total_revenue DESC  
LIMIT 3;
```

### ☛ Explanation:

- Aggregates transaction amounts per region.
- Orders regions by total revenue.
- Retrieves top 3 using LIMIT.

Optional: You could also calculate profit by subtracting costs (if a cost column is present).

## 23. How would you analyze customer churn using SQL?

### ■ Step-by-step SQL approach:

#### ◆ Step 1: Define churn

Let's say a churned customer is one who hasn't transacted in the **last 6 months**.

#### ◆ Step 2: Sample schema

- customers(customer\_id, name, signup\_date)
- transactions(customer\_id, transaction\_date, amount)

---

### ◆ Step 3: Query to identify churned customers

```
SELECT c.customer_id, c.name  
FROM customers c  
LEFT JOIN transactions t  
ON c.customer_id = t.customer_id  
AND t.transaction_date >= CURRENT_DATE - INTERVAL '6 months'  
WHERE t.transaction_id IS NULL;
```

---

### ◆ Step 4: Analyze churn metrics

You could extend this analysis by calculating:

- Churn rate = (Churned Customers / Total Customers) \* 100
  - Monthly churn trend
  - Compare churned vs. active customers in terms of average spend
- 

## 24. Explain the difference between OLAP and OLTP databases.

Feature	OLTP (Online Transaction Processing)	OLAP (Online Analytical Processing)
Purpose	Handles real-time transactional queries	Used for analytical/reporting queries
Operations	INSERT, UPDATE, DELETE	SELECT (aggregate, group, slice, dice)
Data Structure	Highly normalized (3NF)	De-normalized (star/snowflake schema)
Speed	Fast for read/write of single rows	Fast for complex analytical queries
Examples	Banking systems, e-commerce order processing	Business intelligence, dashboards, sales trends
Users	Clerks, DBAs	Analysts, Data Scientists
Backup/Recovery	Essential and frequent	Less frequent

■ In short:

- **OLTP** = operational, fast, real-time transactions.
- **OLAP** = analytical, slow-changing, historical data.

## 25. How would you determine the Average Revenue Per User (ARPU) from transaction data?

◆ ARPU = Total Revenue / Total Number of Users

■ Assume a transactions table:

(transaction\_id, customer\_id, amount, transaction\_date)

■ SQL Query:

```
SELECT  
SUM(amount) * 1.0 / COUNT(DISTINCT customer_id) AS ARPU
```

```
FROM transactions;
```

⌚ **Explanation:**

- SUM(amount) gets total revenue.
- COUNT(DISTINCT customer\_id) counts unique users.
- Multiply by 1.0 to ensure float division.

You can also compute monthly ARPU by grouping by month.

```
SELECT
```

```
    DATE_TRUNC('month', transaction_date) AS month,  
    SUM(amount) * 1.0 / COUNT(DISTINCT customer_id) AS monthly_arpu
```

```
FROM transactions
```

```
GROUP BY month
```

```
ORDER BY month;
```

---

## 26. Describe a scenario where you would use a LEFT JOIN instead of an INNER JOIN.

■ **Use LEFT JOIN when:**

You want **all records from the left table**, even if there's **no matching record** in the right table.

---

◆ **Real-life Scenario:**

**Question:** List all customers and their transactions — even if they haven't made any.

■ **Query:**

```
SELECT c.customer_id, c.name, t.transaction_id, t.amount  
FROM customers c  
LEFT JOIN transactions t  
ON c.customer_id = t.customer_id;
```

⌚ **Why LEFT JOIN?**

- Shows **all customers**, including those with **no transactions** (returns NULLs for those).
- Using INNER JOIN would exclude customers with zero activity.

---

## 27. Write a query to calculate YoY (Year-over-Year) growth for a set of transactions.

Assume a table named transactions with:

(customer\_id, transaction\_date, amount)

■ **Step 1: Extract year-wise revenue**

```
SELECT
```

```
    EXTRACT(YEAR FROM transaction_date) AS year,  
    SUM(amount) AS total_revenue
```

```
FROM transactions
```

```
GROUP BY EXTRACT(YEAR FROM transaction_date);
```

■ **Step 2: Calculate YoY Growth using a CTE and Self-Join**

```
WITH yearly_revenue AS (
```

```

SELECT
    EXTRACT(YEAR FROM transaction_date) AS year,
    SUM(amount) AS total_revenue
FROM transactions
GROUP BY EXTRACT(YEAR FROM transaction_date)
)
SELECT
curr.year AS current_year,
curr.total_revenue,
prev.total_revenue AS previous_year_revenue,
ROUND(((curr.total_revenue - prev.total_revenue) / prev.total_revenue) * 100, 2) AS
yoY_growth_percent
FROM yearly_revenue curr
LEFT JOIN yearly_revenue prev
    ON curr.year = prev.year + 1;

```

**Explanation:**

- Joins each year to its previous year.
- Computes YoY growth as a percentage.

## 28. How would you implement fraud detection using transactional data?

Fraud detection typically involves pattern recognition, anomaly detection, and rule-based filtering.

**Possible SQL-Based Checks:**

Type	Rule
◆ Unusual Amounts	Flag transactions > 3x average amount of that user
◆ Rapid Repeats	Detect multiple transactions from same user within seconds
◆ Location Mismatch	Transactions from different countries within a short time
◆ Card Sharing	Same card used by different customers or IPs

**Example Query – Unusual high amount per user:**

```

WITH avg_txn AS (
    SELECT customer_id, AVG(amount) AS avg_amount
    FROM transactions
    GROUP BY customer_id
)
SELECT t.*
FROM transactions t
JOIN avg_txn a
    ON t.customer_id = a.customer_id
WHERE t.amount > 3 * a.avg_amount;

```

## 29. Write a query to find customers who have used more than 2 credit cards for transactions in a given month.

Assume a transactions table:  
(customer\_id, card\_id, transaction\_date)

■ **Query:**  
SELECT customer\_id,  
 TO\_CHAR(transaction\_date, 'YYYY-MM') AS txn\_month,  
 COUNT(DISTINCT card\_id) AS cards\_used  
FROM transactions  
GROUP BY customer\_id, TO\_CHAR(transaction\_date, 'YYYY-MM')  
HAVING COUNT(DISTINCT card\_id) > 2;

Q **Explanation:**

- Groups by customer\_id and month.
- Counts distinct card\_id used.
- Filters where more than 2 cards were used in a month.

## 30. How would you approach a business problem where you need to analyze the spending patterns of premium customers?

■ **Step-by-Step Structured Approach:**

### ◆ Step 1: Understand the Objective

- Clarify with stakeholders what "**spending pattern**" means.
  - Is it frequency, amount, category, channel, or timing?
- Define "**premium customer**":
  - Based on credit score, card tier (e.g., Platinum, Centurion), monthly spend threshold, etc.

### ◆ Step 2: Data Collection

- Gather relevant datasets:
  - Customer table (ID, tier, demographics)
  - Transactions table (amount, date, category, location)
  - Cards table (card\_type, limits, activation)

### ◆ Step 3: Data Cleaning & Preparation

- Handle missing values and outliers.
- Filter only **premium customers** using defined criteria.
- Enrich data (e.g., categorize merchant types or locations).

### ◆ Step 4: Exploratory Data Analysis (EDA)

Use SQL/Python/Power BI to derive insights like:

Focus Area	Example Analysis
Spend Amount	Average monthly/yearly spend
Time Trends	Seasonality or weekly spending behavior
Categories	Where they spend most (Travel, Dining, Shopping)
Geography	City or region-wise behavior

Focus Area	Example Analysis
Trends	Is their spend increasing/decreasing YoY?

---

### ◆ Step 5: Segmentation

- Use clustering or thresholds to group premium customers into:
  - High spenders
  - Frequent spenders
  - Category loyalists (e.g., only travel)
- Identify anomalies or subgroups with unique patterns.

### ◆ Step 6: Business Recommendations

- Personalize rewards or offers based on their dominant categories.
- Enhance retention strategies for segments showing decline.
- Promote premium card upgrades based on usage patterns.

### ■ Bonus: Sample SQL Query

Get top 3 spending categories of premium customers monthly:

```

SELECT customer_id,
       DATE_TRUNC('month', transaction_date) AS txn_month,
       category,
       SUM(amount) AS total_spend
FROM transactions
WHERE customer_id IN (
    SELECT customer_id FROM customers WHERE tier = 'Premium'
)
GROUP BY customer_id, txn_month, category
ORDER BY customer_id, txn_month, total_spend DESC;
```

---